

```
# Name:Atharva Kangralkar
# SY AIML
# Roll no: 54
# PRN: 12311493
# Colab Link: https://colab.research.google.com/drive/1P_Omhk-BsT6Dhdf0BKC90bleQgeGTxvJ?usp=sharing
```

```
# Q1) Import necessary libraries.
import pandas as pd
import numpy as np
import seaborn as sns
import sklearn as sk
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, classification_report
import matplotlib.pyplot as plt
```

```
# Q2) Load the dataset.
df = pd.read_csv("/content/Social_Network_Ads.csv")
df
```

	User ID	Gender	Age	EstimatedSalary	Purchased	
0	15624510	Male	19	19000	0	
1	15810944	Male	35	20000	0	
2	15668575	Female	26	43000	0	
3	15603246	Female	27	57000	0	
4	15804002	Male	19	76000	0	
...	...	...	...	...	...	
395	15691863	Female	46	41000	1	
396	15706071	Male	51	23000	1	
397	15654296	Female	50	20000	1	
398	15755018	Male	36	33000	0	
399	15594041	Female	49	36000	1	

400 rows × 5 columns

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   User ID         400 non-null   int64
1   Gender          400 non-null   object
2   Age             400 non-null   int64
3   EstimatedSalary 400 non-null   int64
4   Purchased       400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

```
df.describe()
```

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000

```
df = df.drop(columns=["User ID"])
df
```

	Gender	Age	EstimatedSalary	Purchased
0	Male	19	19000	0
1	Male	35	20000	0
2	Female	26	43000	0
3	Female	27	57000	0
4	Male	19	76000	0
...	...	...	...	...
395	Female	46	41000	1
396	Male	51	23000	1
397	Female	50	20000	1
398	Male	36	33000	0
399	Female	49	36000	1

400 rows × 4 columns

Next steps:

Generate code with df

View recommended plots

New interactive sheet

```
df = pd.get_dummies(df)
df = df.astype(int)
```

df

	Age	EstimatedSalary	Purchased	Gender_Female	Gender_Male
0	19	19000	0	0	1
1	35	20000	0	0	1
2	26	43000	0	1	0
3	27	57000	0	1	0
4	19	76000	0	0	1
...	...	...	...	...	...
395	46	41000	1	1	0
396	51	23000	1	0	1
397	50	20000	1	1	0
398	36	33000	0	0	1
399	49	36000	1	1	0

400 rows × 5 columns

Next steps:

Generate code with df

View recommended plots

New interactive sheet

```
# Q3) Select features and target variable.
```

```
X = df.drop(columns=["Purchased"])
# X = df.drop(columns=["Purchased", 'Gender_Female', 'Gender_Male', 'Age'])
y = df["Purchased"]

# Q4) Splitting the dataset into the Training set and Test set.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

X\_train

	Age	EstimatedSalary	Gender_Female	Gender_Male
247	57	122000	1	0
110	39	71000	1	0
16	47	25000	0	1
66	24	19000	0	1
153	36	50000	1	0
...	...	...	...	...
71	24	27000	1	0
106	26	35000	1	0
270	43	133000	1	0
348	39	77000	0	1
102	32	86000	1	0

300 rows × 4 columns

Next steps:

[Generate code with X\\_train](#)

[View recommended plots](#)


[New interactive sheet](#)

y\_train

	Purchased
247	1
110	0
16	1
66	0
153	0
...	...
71	0
106	0
270	0
348	0
102	0

300 rows × 1 columns

```
model = LogisticRegression()
model.fit(X_train, y_train)
```

 /usr/local/lib/python3.11/dist-packages/sklearn/linear\_model/\_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1): STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

▼ LogisticRegression ⓘ ?  
LogisticRegression()

```
accuracy_score(y_test, model.predict(X_test))
#without scaling
```

↩ 0.88

```
# Q4) Splitting the dataset into the Training set and Test set.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

```
# Q5) Perform Feature Scaling. (Use StandardScaler())
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
# Q6) Fit Logistic Regression to the Training set.
model_2 = LogisticRegression()
model_2.fit(X_train, y_train)
```

↩

LogisticRegression ⓘ ?  
 LogisticRegression()

```
accuracy_score(y_test, model_2.predict(X_test))
#after scaling
```

↩ 0.88

```
# Q7) Predict the Test set result.
y_pred = model_2.predict(X_test)
```

```
# Q8) Evaluating the Model. Find out accuracy, confusion matrix, precision, recall, classification report
cm = confusion_matrix(y_test, y_pred)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
```

```
print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}")
print("Classification Report:\n", classification_report(y_test, y_pred))
```

↩


```
Accuracy: 0.88
Precision: 0.9310344827586207
Recall: 0.7297297297297297
Classification Report:
```

	precision	recall	f1-score	support
0	0.86	0.97	0.91	63
1	0.93	0.73	0.82	37
accuracy			0.88	100
macro avg	0.90	0.85	0.86	100
weighted avg	0.89	0.88	0.88	100

```
# Q9) Visualize Confusion Matrix using Heatmap.
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



Confusion Matrix



```
# Q10) Plot Scatterplot of the data points.
plt.figure(figsize=(8, 6))
sns.scatterplot(x=df["Age"], y=df["EstimatedSalary"], hue=y)
plt.xlabel("Age")
plt.ylabel("Estimated Salary")
plt.title("Scatter Plot of Data Points")
plt.show()
```



Scatter Plot of Data Points

